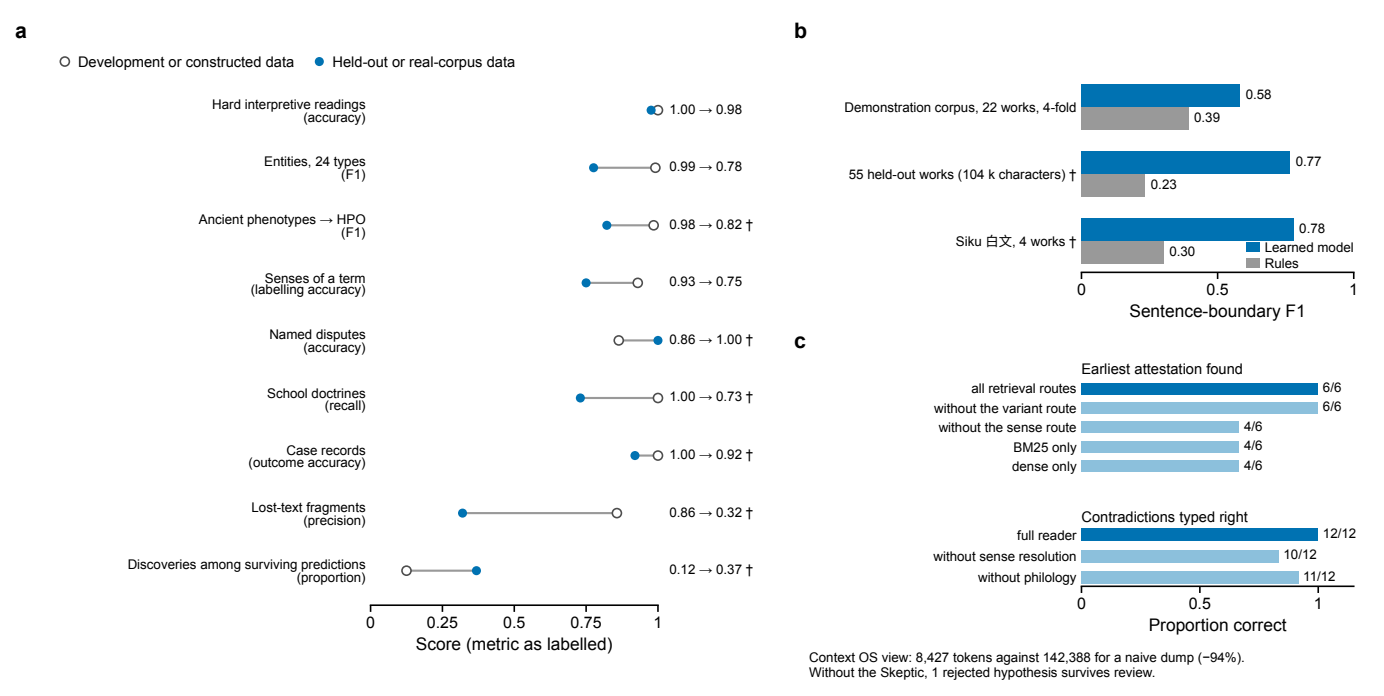
a
Development or constructed data
Held-out or real-corpus data
Hard interpretive readings (accuracy)
1.00 → 0.98
Entities, 24 types (F1)
0.99 → 0.78
Ancient phenotypes → HPO (F1)
0.98 → 0.82 †
Senses of a term (labelling accuracy)
0.93 → 0.75
Named disputes (accuracy)
0.86 → 1.00 †
School doctrines (recall)
1.00 → 0.73 †
Case records (outcome accuracy)
1.00 → 0.92 †
Lost-text fragments (precision)
0.86 → 0.32 †
Discoveries among surviving predictions (proportion)
0.12 → 0.37 †
0
0.25
0.5
0.75
1
Score (metric as labelled)
b
Demonstration corpus, 22 works, 4-fold
0.58
0.39
55 held-out works (104 k characters) †
0.77
0.23
Siku 白文, 4 works †
0.78
0.30
Learned model
Rules
0
0.5
1
Sentence-boundary F1
c
Earliest attestation found
all retrieval routes
6/6
without the variant route
6/6
without the sense route
4/6
BM25 only
4/6
dense only
4/6
Contradictions typed right
full reader
12/12
without sense resolution
10/12
without philology
11/12
0
0.5
1
Proportion correct
Context OS view: 8,427 tokens against 142,388 for a naive dump (−94%).
Without the Skeptic, 1 rejected hypothesis survives review.